

# Multiprocessor Systems

Arquitecturas Paralelas de Computadoras

**Profr. Carlos Ernesto Carrillo Arellano<sup>1</sup>**

<sup>1</sup>Universidad Autónoma Metropolitana - Unidad Azcapotzalco  
Departamento de Ingeniería Electrónica  
Correo electrónico: [ceca@xanum.uam.mx](mailto:ceca@xanum.uam.mx)  
Personal web page: <http://ecarrillo.ddns.net>

**Mayo, 2019**

## 1 Multiprocessor systems

- Introduction
- Parallel-programming model abstractions
- Work partitioning and coordination
- Synchronization and Coordination
- Sequential semantics
- Data Level Parallelism
- Task Level Parallelism

## 2 Shared Memory Systems

## 3 Message Passing System

- Introduction
- Message-passing protocols



## 1 Multiprocessor systems

- Introduction
- Parallel-programming model abstractions
- Work partitioning and coordination
- Synchronization and Coordination
- Sequential semantics
- Data Level Parallelism
- Task Level Parallelism

## 2 Shared Memory Systems

## 3 Message Passing System

- Introduction
- Message-passing protocols



# Introduction

- 1 Since the beginning of the history of computer systems, the demand for **more performance** has been the most important driving force for evolution in computer architecture
- 2 Many important **applications** demand more performance than a single (serial) processor core can provide, and historically have pushed parallel architecture technology
- 3 **Examples:** Numerical programs used in computer simulation to analyze and solve problems in science and engineering, such as climate modeling, weather forecasting, or computer-aided design



# Introduction

- 1 ¿One processor form many tasks?
- 2 ¿One processor for an intensive task?
- 3 Interest in **multiprocessor** architecture research is focus os increasing the performance of computer systems.
- 4 A multiprocessor is a computer system with **two or more central processing units** (CPUs) **linked together** to enable parallel processing
- 5 **Cooperation and communication**
- 6 The key objective of using a multiprocessor is to boost the **system's execution speed**, with other objectives being fault tolerance and multitasking




# Introduction

- 1 The main reason for parallel programming is to **execute code efficiently**, since parallel programming saves time, allowing the **execution of applications in a shorter wall-clock time**.
- 2 As a consequence of executing code efficiently, **parallel programming often scales with the problem size, and thus can solve larger problems**.
- 3 Parallel programming is a means of providing concurrency, particularly performing **simultaneously multiple actions at the same time**.




# Shared-memory and message-passing

- 1 There are two multiprocessor architectural styles
- 2 **Shared-memory and message-passing** multiprocessor systems
- 3 Both styles use multiple processors with the goal of achieving a **linear speedup of computational power with the number of processors.**
- 4 They differ in the method by which the **processors exchange data.**



# Shared-memory and message-passing


- 1 Processors in shared-memory multiprocessors share the same address space and can exchange data through **shared-memory locations by regular load and store instructions.**
- 2 Processors in message-passing multiprocessor systems have their own (private) **address space** and communicate data between their address spaces by exchanging **messages**
- 3 A fundamental aspect is that software must be written in such a way that parallelism is exposed





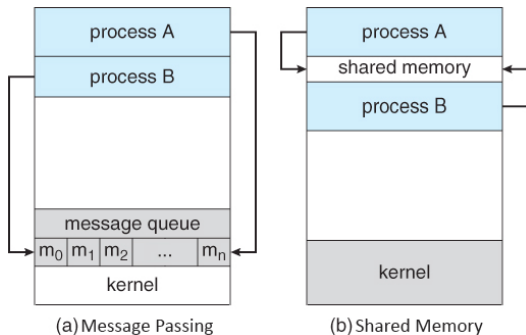
# Shared-memory and message-passing

- 1 Computational tasks that can run in parallel must first be **identified and partitioned across the processors in a balanced fashion**
- 2 **Dependences** between tasks and **communication** needed to transfer results among tasks, **collectively called coordination**, must be taken into deep consideration
- 3 **Primitives for synchronization and communication** constitute an important part of what the underlying architecture has to support and expose to the software





# Shared-memory and message-passing

- 1 In shared-memory multiprocessor systems, communication is intuitively supported by the inherent shared-memory address space offered by these systems, although **explicit support for synchronization is needed**
- 2 Message-passing systems, on the other hand, offer explicit primitives for synchronization as well as communication: **send and receive primitives**




# Parallel-programming model



- 1 A parallel-programming model defines how parallel computations can be expressed in a **high level programming language**
- 2 Popular parallel programming models have been implemented as simple **extensions** to commonly used programming languages
- 3 **Message Passing Interface** (MPI) for message-passing
- 4 **OpenMP** for shared-memory systems

# Parallel-programming model

- 1 Regardless of the parallel computer architecture-style (shared-memory or message-passing) **programmers or compilers must be able to express parallelism**
- 2 **Work partitioning** and **coordination**
- 3 Work that can be carried out in parallel must be **identified and partitioned among the processors**
- 4 **Thread** or **process** interchangeably: code that is run on a single processor (or core) in a parallel computer






## Parallel-programming model

- 1 **Work** done in parallel by threads running on different processors must be **coordinated** so that **the end result is the same as if the work was done by a single processor**.
- 2 Coordination involves two actions
- 3 One to **synchronize** parallel threads
- 4 **Communicate** partial results between threads
- 5 **Threads need to exchange information** either through the **memory**, or by **explicit messaging**

# Synchronization and Coordination

- 1 Regardless of the communication model, **communication** obviously **takes time** and has an impact on how fast a problem can be solved on a parallel computer
- 2 From an architecture point of view, it is important to provide adequate **support for synchronization and communication** so that they can be carried out efficiently.
- 3 Remove bottlenecks in a shared-memory system or in the interconnection network.




# Speedup

- 1 A key reference point for both the architect and the application developer is how the use of parallelism **improves the performance of the application**.
- 2 A typical measure is the improvement in **execution time**
- 3 For a single, fixed problem, the performance of the machine on the problem is simply the **reciprocal of the time to complete the problem**


$$\text{Speedup}_{\text{fixed problem}}(p \text{ processors}) = \frac{\text{Time}(1 \text{ processor})}{\text{Time}(p \text{ processors})}.$$

# Shared memory vs message passing





# Shared memory vs message passing



**Who does communication?**  
**Data distribution**  
**HW Support**  
**Programmer**  
**Correctness**  
**Performance**

### **Message Passing**

Programmer  
Manual  
Simple (NIC)

Difficult  
Difícil


### **Shared Memory**

Automatic  
Automatic  
Extensive

Less Difficult  
Very Difficult


# Parallel-programming

- 1 **Starting** from a sequential program, the **programmer or compiler** must first identify the **parts of the program which can be run in parallel**.
- 2  $S_1$  and  $S_2$  are two **program segments** that are executed one after the other in the sequential program
- 3 Can run in parallel if, and only if,  $S_1$  is **independent** of  $S_2$ , meaning that  $S_1$  does not produce data used by  $S_2$
- 4 Running  $S_1$  and  $S_2$  in parallel yields the same result as if they **were executed one after the other**
- 5 The parallel program then conforms to **sequential semantics**




# Parallel-programming

- 1 Finding independent code segments in a program is **key to unlocking the parallelism** exploitable by multiprocessors
- 2 **Data-level** parallelism means that computations of different data elements are **independent of each other**.
- 3 **Data-level** parallelism is often expressed in **loops** and is one of the main targets for parallelizing compilers to unlock parallelism
- 4 To exploit data-level parallelism in loops, it must make **sure that there is no loop-carried dependency**, meaning that the computations in different iterations are independent of each other.



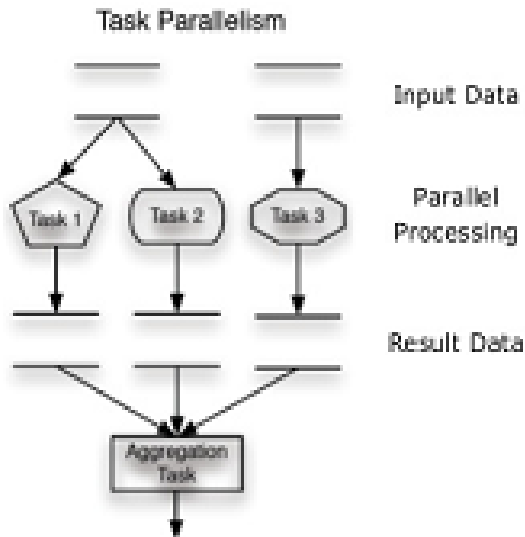
# Data level parallelism

- 1 When the **same computation is applied to all data elements** in an array, the term SPMD parallelism is often used
- 2 **SPMD** stands for **single-program-multiple-data**
- 3 The **same function (program code)** is **applied to all data elements**
- 4 Data-level parallelism has the attractive property that **as the problem size is scaled up, more parallelism can often be found**.



## Task Level Parallelism

- 1 In this form of parallelism **independent functions are executed on different processors**.
- 2 These functions can then form a **function pipeline** in which the data stream is processed by applying a **sequence of functions** to it.
- 3 **Data-level parallelism is applied to each function of a function pipeline to exploit parallelism at two levels.**



## 1 Multiprocessor systems

- Introduction
- Parallel-programming model abstractions
- Work partitioning and coordination
- Synchronization and Coordination
- Sequential semantics
- Data Level Parallelism
- Task Level Parallelism

## 2 Shared Memory Systems

## 3 Message Passing System

- Introduction
- Message-passing protocols



# Problem

```
1 sum = 0;
2 for (i=0,i<N,i++)
3     for (j=0,j<N,j++){
4         C[i,j] = 0;
5         for (k=0,k<N,k++)
6             C[i,j] = C[i,j] + A[i,k]*B[k,j];
7         sum += C[i,j];
8     }
```

- 1 Multiplication of two matrices **A and B**
- 2 Each with **N rows and N** columns
- 3 The **sum of all matrix elements**
- 4 The result C is a  $N \times N$  matrix and a scalar variable `sum` stores the sum of all matrix elements
- 5 This program has lots of **data-level parallelism** because the calculation of each individual matrix element is independent of the calculation of others.
- 6 All matrix elements could be **calculated in parallel**.



# Problem

```
1 sum = 0;
2 for (i=0,i<N,i++)
3     for (j=0,j<N,j++){
4         C[i,j] = 0;
5         for (k=0,k<N,k++)
6             C[i,j] = C[i,j] + A[i,k]*B[k,j];
7         sum += C[i,j];
8     }
```

- 1 Multiplication of two matrices **A and B**
- 2 Each with **N rows and N** columns
- 3 The **sum of all matrix elements**
- 4 The result C is a  $N \times N$  matrix and a scalar variable `sum` stores the sum of all matrix elements
- 5 This program has lots of **data-level parallelism** because the calculation of each individual matrix element is independent of the calculation of others.
- 6 All matrix elements could be **calculated in parallel**.

# Problem

```
1 sum = 0;
2 for (i=0,i<N,i++)
3     for (j=0,j<N,j++){
4         C[i,j] = 0;
5         for (k=0,k<N,k++)
6             C[i,j] = C[i,j] + A[i,k]*B[k,j];
7         sum += C[i,j];
8     }
```

- 1 Multiplication of two matrices **A and B**
- 2 Each with **N rows and N** columns
- 3 The **sum of all matrix elements**
- 4 The result C is a  $N \times N$  matrix and a scalar variable `sum` stores the sum of all matrix elements
- 5 This program has lots of **data-level parallelism** because the calculation of each individual matrix element is independent of the calculation of others.
- 6 All matrix elements could be **calculated in parallel**.

# Shared Memory Systems

- 1 The **parallelism granularity** is the **amount of work that is carried out by a thread in parallel with other threads**.
- 2 For a fixed problem size, for example a matrix with  $N \times N$  elements, **the granularity decreases with the number of processors**.
- 3 Under a shared-memory programming model, **variables can be defined globally as shared**.
- 4 This has the attractive property that the source data **can be accessed by all threads**.
- 5 Each thread can **read** from the input matrices A and B and **store** its results into matrix C using regular loads and stores



# Shared Memory Systems

- 1 The **parallelism granularity** is the **amount of work that is carried out by a thread in parallel with other threads.**
- 2 For a fixed problem size, for example a matrix with  $N \times N$  elements, **the granularity decreases with the number of processors.**
- 3 Under a shared-memory programming model, **variables can be defined globally as shared.**
- 4 This has the attractive property that the source data **can be accessed by all threads.**
- 5 Each thread can **read** from the input matrices A and B and **store** its results into matrix C using regular loads and stores



# Shared Memory Systems

- 1 The **parallelism granularity** is the **amount of work that is carried out by a thread in parallel with other threads.**
- 2 For a fixed problem size, for example a matrix with  $N \times N$  elements, **the granularity decreases with the number of processors.**
- 3 Under a shared-memory programming model, **variables can be defined globally as shared.**
- 4 This has the attractive property that the source data **can be accessed by all threads.**
- 5 Each thread can **read** from the input matrices A and B and **store** its results into matrix C using regular loads and stores

# Shared Memory Systems

- 1 The **parallelism granularity** is the **amount of work that is carried out by a thread in parallel with other threads.**
- 2 For a fixed problem size, for example a matrix with  $N \times N$  elements, **the granularity decreases with the number of processors.**
- 3 Under a shared-memory programming model, **variables can be defined globally as shared.**
- 4 This has the attractive property that the source data **can be accessed by all threads.**
- 5 Each thread can **read** from the input matrices A and B and **store** its results into matrix C using regular loads and stores

# Shared Memory Systems

- 1 The **parallelism granularity** is the **amount of work that is carried out by a thread in parallel with other threads.**
- 2 For a fixed problem size, for example a matrix with  $N \times N$  elements, **the granularity decreases with the number of processors.**
- 3 Under a shared-memory programming model, **variables can be defined globally as shared.**
- 4 This has the attractive property that the source data **can be accessed by all threads.**
- 5 Each thread can **read** from the input matrices A and B and **store** its results into matrix C using regular loads and stores

# Shared Memory Systems

- 1 Each thread calculates its partial sum and **the partial sums must be accumulated in the global sum**
- 2 Since there is a **risk** that **multiple threads could each read the global sum into a processor register**, add their **partial sum** to it, and then **write back the result into the global sum variable**, some of the **partial sums could be overwritten by others**.
- 3 Therefore, the partial sums must be added to the **global sum in a serial fashion**, inside a **critical section**.



# Shared Memory Systems

- 1 Each thread calculates its partial sum and **the partial sums must be accumulated in the global sum**
- 2 Since there is a **risk** that **multiple threads could each read the global sum into a processor register**, add their **partial sum** to it, and then **write back the result into the global sum variable**, some of the **partial sums could be overwritten by others**.
- 3 Therefore, the partial sums must be added to the **global sum in a serial fashion**, inside a **critical section**.

# Shared Memory Systems

- 1 Each thread calculates its partial sum and **the partial sums must be accumulated in the global sum**
- 2 Since there is a **risk** that **multiple threads could each read the global sum into a processor register**, add their **partial sum** to it, and then **write back the result into the global sum variable**, some of the **partial sums could be overwritten by others**.
- 3 Therefore, the partial sums must be added to the **global sum in a serial fashion**, inside a **critical section**.

# Shared Memory Systems

- 1 The semantics of a critical section ensures that **at most one thread can execute the code inside of the section at a time.**
- 2 All threads add their partial sums to the global sum one at a time so that the end result is the correct sum of all partial sums.
- 3 Since threads run **asynchronously**, it may happen that one thread is finished with its partial matrix product before another one has even started.
- 4 No thread can enter the critical section and update sum until all threads are done with their part of the matrix multiplication algorithm
- 5 **Barrier synchronization** forces every thread to stall until all threads have reached the barrier.



# Shared Memory Systems

- 1 The semantics of a critical section ensures that **at most one thread can execute the code inside of the section at a time.**
- 2 All threads add their partial sums to the global sum one at a time so that the end result is the correct sum of all partial sums.
- 3 Since threads run **asynchronously**, it may happen that one thread is finished with its partial matrix product before another one has even started.
- 4 No thread can enter the critical section and update sum until all threads are done with their part of the matrix multiplication algorithm
- 5 **Barrier synchronization** forces every thread to stall until all threads have reached the barrier.



# Shared Memory Systems

- 1 The semantics of a critical section ensures that **at most one thread can execute the code inside of the section at a time.**
- 2 All threads add their partial sums to the global sum one at a time so that the end result is the correct sum of all partial sums.
- 3 Since threads run **asynchronously**, it may happen that one thread is finished with its partial matrix product before another one has even started.
- 4 No thread can enter the critical section and update sum until all threads are done with their part of the matrix multiplication algorithm
- 5 **Barrier synchronization** forces every thread to stall until all threads have reached the barrier.



# Shared Memory Systems

- 1 The semantics of a critical section ensures that **at most one thread can execute the code inside of the section at a time.**
- 2 All threads add their partial sums to the global sum one at a time so that the end result is the correct sum of all partial sums.
- 3 Since threads run **asynchronously**, it may happen that one thread is finished with its partial matrix product before another one has even started.
- 4 No thread can enter the critical section and update sum until all threads are done with their part of the matrix multiplication algorithm
- 5 **Barrier synchronization** forces every thread to stall until all threads have reached the barrier.



# Shared Memory Systems

- 1 The semantics of a critical section ensures that **at most one thread can execute the code inside of the section at a time.**
- 2 All threads add their partial sums to the global sum one at a time so that the end result is the correct sum of all partial sums.
- 3 Since threads run **asynchronously**, it may happen that one thread is finished with its partial matrix product before another one has even started.
- 4 No thread can enter the critical section and update sum until all threads are done with their part of the matrix multiplication algorithm
- 5 **Barrier synchronization** forces every thread to stall until all threads have reached the barrier.



# Shared Memory Systems

- 1 Shared-memory multiprocessors have the attractive property that communication and synchronization can be carried out by primitives available in the ISA
- 2 Loads and stores may sometimes suffer from very long access latencies
- 3 The design of the underlying memory system is critical to performance.

```
/* A, B, C, BAR, LV and sum are shared
/* All other variables are private
1a low = pid*N/nproc; /* pid=0...nproc-1
1b hi = low + N/nproc; /* identifies rows of A
1c mysum = 0; sum = 0;
2 for (i=low,i<hi,i++)
3     for (j=0,j<N,j++){
4         C[i,j] = 0;
5         for (k=0,k<N,k++){
6             C[i,j] = C[i,j] + A[i,k]*B[k,j];
7             mysum +=C[i,j];
8         }
9 BARRIER(BAR);
10 LOCK(LV);
11     sum += mysum;
12 UNLOCK(LV);
```



## 1 Multiprocessor systems

- Introduction
- Parallel-programming model abstractions
- Work partitioning and coordination
- Synchronization and Coordination
- Sequential semantics
- Data Level Parallelism
- Task Level Parallelism

## 2 Shared Memory Systems

## 3 Message Passing System


- Introduction
- Message-passing protocols



# Introduction

- 1 In shared-memory systems, where **coordination happens through an address space shared among all threads**
- 2 In the message-passing programming model, each thread or process has **its own address space**.
- 3 **Coordination** is carried out by sending **explicit messages** between the threads.





## Message-passing systems

- 1 Message passing can be supported in any system that consists of a number of **interconnected computational nodes** and where **each node has at least one processor and some memory**.
- 2 Two major implications:
- 3 Data structures must be explicitly **distributed to the private address spaces**
- 4 Results from the partial computations performed by the threads must be **collected at the end**

# Message-passing systems

- 1 Multiplication of matrices A and B
- 2 A and B are initially kept in one computational node, the **master node**
- 3 The master node **partitions** the matrices across the nodes
- 4 **Send** and **Receive**
- 5 **SEND** copies data from the sender's local address space to a buffer at the receiver's side
- 6 **RECV** copies data from that buffer to the local address space at the receiver's side

```
1a myN = N/nproc;
1b if(pid == 0)
1c   for(i=1; i<nproc;i++){
1d     k=i*N/nproc;
1e     SEND(&A[k][0],myN*N*sizeof(float),i,IN1);
1f     SEND(&B[0][0],N*N*sizeof(float),i,IN2);
1g   } else {
1h     RECV(&A[0][0],myN*N*sizeof(float),0,IN1);
1i     RECV(&B[0][0],N*N*sizeof(float),0,IN2);
1j   }
1k mysum = 0;
2   for (i=0,i<myN, i++)
3     for (j=0,j<N, j++){
4       C[i,j] = 0;
5       for (k=0,k<N, k++)
6         C[i,j] = C[i,j] + A[i,k]*B[k,j];
7       mysum += C[i,j];
8     }
9   if (pid == 0){
10    sum = mysum;
11    for(i = 1;i<nproc;i++){
12      RECV(&mysum,sizeof(float),i,SUM);
13      sum += mysum;
14    }
15    for(i=1; i<nproc;i++){
16      k=i*N/nproc;
17      RECV(&C[k][0],myN*N*sizeof(float),i,RES);
18    }
19  } else{
20    SEND(&mysum,sizeof(float),0,SUM);
21    SEND(&C[0][0],myN*N*sizeof(float),0,RES);
22  }
```

# Message-passing systems

- 1 Each process works on the partition assigned
- 2 Local portions of the result matrix must later be copied back to the master node
- 3 In the message-passing system implementation, **synchronization** is **implicit** in the message-passing primitives.
- 4 SENDs and RECVs usually come in two main flavors: **synchronous** and **asynchronous**
- 5 Synchronous SENDs and RECVs **block until both parties have notified each other that the message has been exchanged**

```
1a myN = N/nproc;
1b if(pid == 0)
1c   for(i=1; i<nproc;i++){
1d     k=i*N/nproc;
1e     SEND(&A[k][0],myN*N*sizeof(float),i,IN1);
1f     SEND(&B[0][0],N*N*sizeof(float),i,IN2);
1g   } else {
1h     RECV(&A[0][0],myN*N*sizeof(float),0,IN1);
1i     RECV(&B[0][0],N*N*sizeof(float),0,IN2);
1j   }
1k mysum = 0;
2   for (i=0,i<myN, i++)
3     for (j=0,j<N, j++){
4       C[i,j] = 0;
5       for (k=0,k<N, k++){
6         C[i,j] = C[i,j] + A[i,k]*B[k,j];
7         mysum += C[i,j];
8       }
9   if (pid == 0){
10    sum = mysum;
11    for(i = 1;i<nproc;i++){
12      RECV(&mysum,sizeof(float),i,SUM);
13      sum += mysum;
14    }
15    for(i=1; i<nproc;i++){
16      k=i*N/nproc;
17      RECV(&C[k][0],myN*N*sizeof(float),i,RES);
18    }
19  } else{
20    SEND(&mysum,sizeof(float),0,SUM);
21    SEND(&C[0][0],myN*N*sizeof(float),0,RES);
22  }
```



# Message-passing multiprocessor

- 1 In Message-passing multiprocessor architectures **participating nodes exchange explicit messages** among each other
- 2 Only the nodes involved in a particular exchange are contacted
- 3 Message-passing systems can be easily built on top of a cluster of desktop/laptop machines as well as embedded on top of a shared-memory system.





# Message-passing primitives

- 1 **Message exchanges** are the fundamental primitive for **synchronization** as well as **communication**
- 2 Synchronous message-passing
- 3 A matching pair of SEND and RECV primitives sets up a communication path
- 4 Data can be copied from a designated location in the sender's local address space to a designated location in the receiver's local address space

# Message-passing primitives

- 1 **SEND** (location in the local address space (&A), length of the message (sizeof(A)), ID of the receiver (P2), tag to distinguish the message from other messages sent between P1 and P2)
- 2 **RECV** (location in the receiver's local address space (&B), the size of B (which should match the size in the SEND), the ID of the sender (P1), and the same tag as the matching SEND)
- 3 P1 and P2 run **asynchronously**, the issue for the SEND RECV protocol is to **guarantee** that the content of A will be transferred and correctly copied into B before P2 executes its statement after returning from the RECV, which actually happens to change the content of B.

Code for process P1:

```
A=10;  
SEND (&A, sizeof (A) , P2, SEND_A) ;  
A=A+1;  
RECV (&C, sizeof (C) , P2, SEND_B) ;  
printf (C) ;
```

Code for process P2:

```
B=5;  
RECV (&B, sizeof (B) , P1, SEND_A) ;  
B=B+1;  
SEND (&B, sizeof (B) , P1, SEND_B) ;
```



# Synchronous message passing

- 1 The sender is **block** until the message has been received by the receiver
- 2 The receiver is **block** until the message is available and has been copied into the designated data structure in its local address space.
- 3 There are **two disadvantages** with synchronous message exchanges: one is that they are prone to **deadlock**
- 4 The second is that they do not allow overlapping communication with computation

Code for process P1:

```
A = 10;
SEND (&A, sizeof (A) , P2, SEND_A) ;
RECV (&B, sizeof (B) , P2, SEND_B) ;
```

Code for process P2:

```
B = 5;
SEND (&B, sizeof (B) , P1, SEND_B)
RECV (&A, sizeof (B) , P1, SEND_A) ;
```

# Asynchronous message passing

Code for process P1:

```
A=10;  
SEND (&A, sizeof (A) , P2, SEND_A) ;  
    <UNRELATED COMPUTATION;>  
RECV (&B, sizeof (B) , P2, SEND_B) ;
```

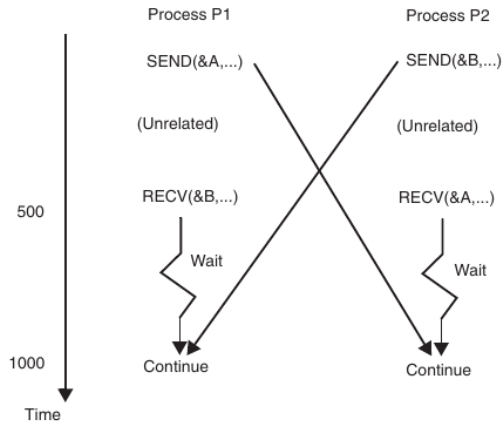
Code for process P2:

```
B=5;  
SEND (&B, sizeof (B) , P1, SEND_B) ;  
    <UNRELATED COMPUTATION;>  
RECV (&A, sizeof (B) , P1, SEND_A) ;
```

- 1 Synchronous messages transfer is that they **combine synchronization with communication**
- 2 **Long-latency operations**, and packing them into a single primitive can lead to **performance losses**
- 3 **Separating** them could allow the sender to **do useful work while the message is in transit**.
- 4 **Asynchronous message passing** primitives do just that

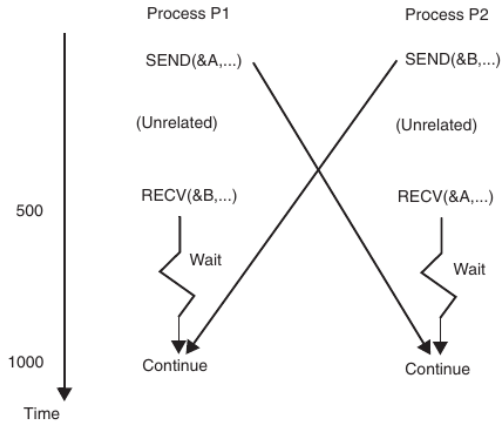
# Asynchronous message passing

- 1 P1 sends the message and then moves on to execute the code after its SEND, **without waiting** for the receiver to copy the message into its address space.
- 2 The code after SEND is, in this case, **unrelated** in the sense that it does not change the content of A, nor does it need the value of B



# Asynchronous message passing

- 1 It could happen that the content of A that is copied does not correspond to what it was at the time the SEND was executed, as the message transfer happens asynchronously with the execution of the code
- 2 **Two forms of asynchronous message-passing primitives exist, blocking and non-blocking message-passing primitives.**
- 3 In **asynchronous message** passing the sender does not necessarily wait until the data in its local address space have been copied.



# Blocking and non-blocking message-passing

- 1 A **blocking asynchronous** SEND gives back control to the sending process once a copy of the local data making up the message has been buffered somewhere and cannot be affected by the execution of the sending process
- 2 A **blocking RECV** does not give back control to the receiving process until after the message has been copied into the local address space of the receiver.
- 3 **Non-blocking asynchronous** message-passing primitives, where the control is returned to the sender and receiver immediately and where the transfer happens in the background






## Probe function

- 1 Some message-passing systems also provide a probe function that can **interrogate the status of the message transfer**
- 2 Probe functions make it possible to **check whether data has been copied from the local address space of the sender to a buffer** or whether it has been copied into the **local address space of the receiver**

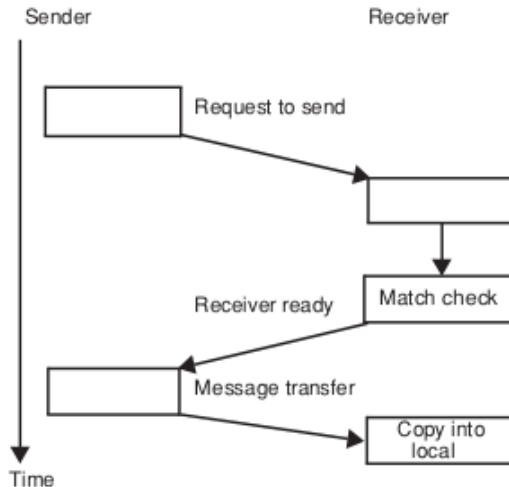
# Introduction

- 1 Consider a **synchronous message transfer**
- 2 When sender executes its SEND, it **needs to synchronize** with the receiver to make sure that it has reached its matching RECV.
- 3 At this point, the message with **the data copied from the local address space of the sender** can be sent off to the receiver, which then **copies it into the location specified in its local address space**



# Message-passing protocol

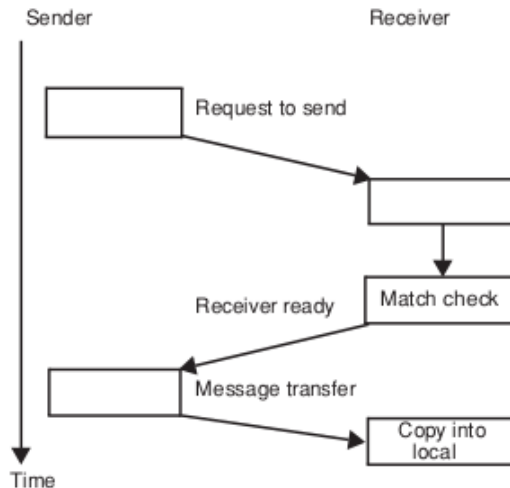
- 1 A protocol is used to implement the message transfer
- 2 When the sender process invokes the SEND function, a **message asking whether the receiver is ready** is sent to the receiver ("Request to send")
- 3 In order for the receiving node to figure out whether it is ready for the message transfer, it **interrogates a table – the match table** – which keeps track of the status of all RECVs





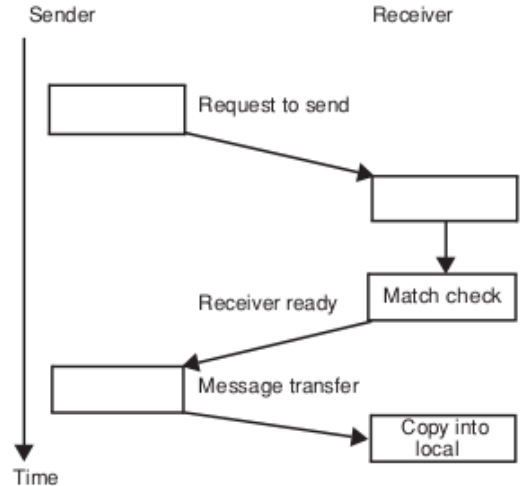
# Message-passing protocol

- 1 There are two cases: (1) **the matching RECV has already been executed**, or (2) **the matching RECV has not yet been executed**.
- 2 If the receiver process has **already executed the matching RECV**, there will be a **match in the table** and the receiving node notifies the sending node that it is ready for the message transfer ("**Receiver ready**")



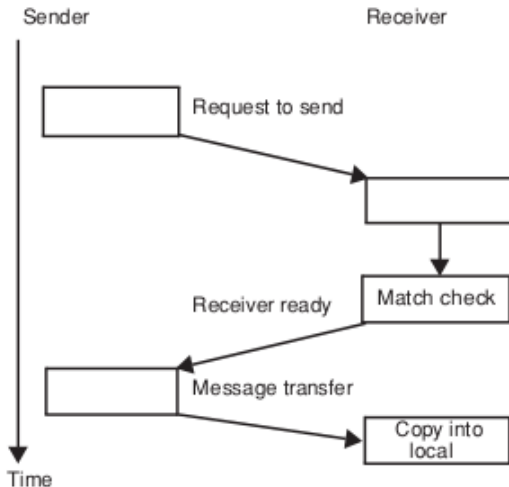
# Message-passing protocol

- 1 There are two cases: **(1) the matching RECV has already been executed**, or **(2) the matching RECV has not yet been executed**.
- 2 If the receiver process has **already executed the matching RECV**, there will be a **match in the table** and the receiving node notifies the sending node that it is ready for the message transfer ("**Receiver ready**")
- 3 The sender then sends the message to the receiver, which copies it into the specified area in its local address space.



# Message-passing protocol

- 1 If the receiving node has not yet executed the RECV function, a **“Receiver ready” message does not go out until the matching RECV function is executed** and the match table is updated accordingly
- 2 A three-phase protocol is needed to send the message.
- 3 Protocol for synchronous message-passing primitives
- 4 This means that the sender and the receiver processes are **both blocked until the message transfer is completed**



# Message-passing protocol



- 1 ¿What happens under blocking asynchronous message passing?
- 2 Under blocking asynchronous message passing, the sender can continue executing past the SEND function **once the message is buffered**
- 3 Buffering space has to be reserved to host a copy of the message data of the sender before the sender can resume its execution past the SEND function
- 4 copy of the content of the local data structure is created in parallel with sending a request to the receiver
- 5 Once the copy is created, the sender can resume its execution past the SEND function.

# Introduction

- 1 Message-passing systems can be built on nodes connected by an interconnection network
- 2 The protocols can be built upon the supported low-level network transaction primitives.
- 3 These primitive network transactions are provided by general interconnection networks.
- 4 Additional hardware support for message transfer aims to cut down the latency of sending a message from one node to another